




The Context Engineering Lifecycle

A 7-Phase Blueprint for Reliable AI

Instead of treating context engineering as a single act of prompt writing, this framework treats it as an operational lifecycle. Move from intuitive guessing to durable, artifact-driven problem solving.

A reliable demo is not a reliable design.




Generative AI work is vulnerable to hidden design debt. When teams focus only on wording the prompt, they fix symptoms while leaving the real design problem untouched.

Moving from intuition to engineering.

Intuitive Prompting		Context Engineering	
One-shot creative exercise	➡	Disciplined operational lifecycle	
Subjective gut-feel evaluation	➡	Evidence-based rubric scoring	
Start by writing instructions	➡	Define success before writing	
Lives in chat history	➡	Produces durable artifacts	
Visible artifacts turn subjective debate into reviewable evidence.			

The 7-Phase Context Engineering Lifecycle



In practice, this is not a one-time march from left to right. It is a controlled loop. Evolving expectations and new policies mean deployment feeds right back into discovery.

Phase 1: Define success before writing prompts.


Translate broad ambitions into a bounded job-to-be-done with explicit success criteria and representative test queries.




Bounded Constraints

- ☰ Target Task
- 👤 Intended Users
- ✓ Success Criteria
- ! Edge Cases

Durable Artifacts Ecosystem




Problem Statement



Test Cases

Phase 2: Filter for high-signal context

The context window is finite. Over-selection competes for the model's attention and dilutes critical instructions.






Durable Artifacts Ecosystem




Phase 3: Build a reviewable context skeleton

Separate instructions, evidence, dynamic facts, and user requests so the model can distinguish rules from data.



Phase 4: Execute on representative reality.

Run the design on ordinary, difficult, and borderline cases. Capture outputs systematically to serve as an evidence baseline.




**Durable
Artifacts
Ecosystem**



Problem
Statement



Test Cases



Approved
Sources



Budget Plan



Run Logs



Failure Patterns


The 0-1-2 Evaluation Scorecard

A repeatable framework to replace subjective gut feel


Criteria	0 (Missed / Violated)	1 (Partial / Inconsistent)	2 (Met Clearly / Reliably)
Accuracy	✗ Materially wrong or unsupported	— Partly right but incomplete or uncertain	✓ Correct and appropriately grounded
Completeness	✗ Misses major needed content	— Addresses part of the task	✓ Covers the task to the needed depth
Constraint Adherence	✗ Violates key rule or boundary	— Mostly compliant but inconsistent	✓ Consistently follows stated rules
Usability	✗ Hard to act on	— Partly usable	✓ Clear, structured, ready to use

Phase 5: Score evidence, not impressions.

Compare captured outputs against explicit success criteria to identify recurring failure modes.




Durable Artifacts Ecosystem



Phase 6: Refine through controlled regression

Adjust one meaningful variable at a time to fix specific failures without breaking previously successful outputs.



Durable Artifacts Ecosystem


Problem
Statement


Test Cases


Approved
Sources


Budget
Plan


Run Logs


Failure
Patterns



Evaluation
Report


Versioned
Template


Iteration
Log

Phase 7: Operationalize and monitor for drift.

Deployment is not the end. Define ownership, monitor logging, and establish an update policy as software or user queries change.



Durable Artifacts Ecosystem




The cost of skipping discipline.

Each skipped phase removes a type of discipline that later phases depend on.

The Skipped Phase	Predictable Failure Pattern
Skipped Discovery	Optimizing the wrong task because success was never defined precisely.
Skipped Selection	Context overload, causing the model to become verbose and distracted.
Skipped Shaping	Mixed prompts where rules and facts blur, making failures hard to diagnose.
Skipped Evaluation	Random rewrites and guesswork replacing targeted iteration.

Linear is an illusion: The Controlled Loop.

The true output of this lifecycle isn't just a deployed AI model; it's an adaptable, self-correcting business system.



	Criteria		Score	Reward	Scoring	Prompt Template:			Prompt
Criteria	- Evaluates content and relevance to the user's request. - Ensures the response is accurate and up-to-date. - Keeps the response concise and to the point.	1	1	0		<pre>System: You are a helpful assistant. User: [User's question] Assistant: [Your response]</pre>			
Deployment	- Identify the target audience and their needs. - Develop a content strategy that aligns with the brand's goals. - Choose the right channels to reach the audience.	2	0	3					
Prompt Templates	- Identify the user's intent and the information they need. - Use the information to generate a response that is both helpful and concise.	2	3	4					
Backend Processes	- Ensure the system is scalable and can handle a large number of users. - Implement security measures to protect user data.	1	7	2					

The Knowledge Worker's Advantage.

You own the business problem. Context engineering provides, provides the operational discipline to ensure generative AI solves it reliably.

Budget Planning			
Planary	Rrice	Scerting	Reasbes
Budget	\$96.00	\$2.060	\$2.00
Budget Planning	\$100	\$4.00	\$4.00
Budget Planning	\$20.00	-\$2.00	\$7.00
Total / Total	\$26.00	\$0.30	\$6.00

