


Regression Watchtower

A Visual Guide to Baseline Protection and
Evidence-Based Release Judgment



Chapter 4 • Reliability, Failure Handling, and Measurement (Build v15.0.0)

The Baseline-First Rule



The Baseline



The Distraction

Before judging if a candidate sounds smoother, identify:


What problem was the baseline solving?

What evidence did it keep visible?

What boundaries did it enforce?

Do not reward the candidate for tone before naming what the baseline protected

The 5-Phase Learning Loop



↺ Replay for Mastery

Use the 'Try mission again' button to practice the reasoning pattern until you can explain the decision without teaching aids.

Phase 1: Launcher & Mission Selection

CHAPTER 4 • RELIABILITY, FUTURE HAZARDING, AND MEASUREMENT

Regression Watchtower

A learning app for knowledge workers practicing boring protection, regression awareness, and evidence-based release judgment.

Version v1.0.0 | Environment Local | Occasion Not started

Use dark theme | Glossary | Help | Feedback

MISSION LAUNCHER

Choose a watchtower run

Each mission is built around the same Chapter 4 ditherer solvent to the incident, anchor on the baseline, compare what the baseline protects, choose a choice judgment, and decide what skill needs to be watched after the change.

5-10 minute learning run
8 authored scenarios

Available incidents

Use filters to choose a domain, difficulty, pack, or completion mode. Then work one mission from incident review to adocot.

SEARCH | DOMAIN | DIFFICULTY | PACK | MODE | SORT

Find by MLC, abjotix | All done | All level | All pack | All mod | Recent | Best Higgs

6 missions currently match these filters.

Compliance / Support | Foundation | Internal only completion | Core Watchtower Pack

Support Summary Compression Update

Decide whether a shorter support summary format preserved the baseline's acclusability while refusing noise.

Learning term: Protest the trusted baseline, judge the candidate change, and choose monitoring proportionate to the remaining risk.

Start mission

ML / Hilly | Foundation | External completion step | Core Watchtower Pack

Travel Policy Qiskit Prompt Change

Judge whether a friendlier HRT answer style preserved the baseline's policy grounding and occlusion discipline.

Learning term: Protect the model baseline, judge the candidate change, and choose monitoring proportionate to the remaining risk.

Start mission

Compliance / Openness (tagging) | Intermediate | Internal only completion | Core Watchtower Pack

Executive Digest Simplification Drift

Decide whether a shorter routine digest preserved the baseline's warning signals instead of hiding them before a more positive headline.

Start mission

What you are producing

- Protest a known good baseline before accepting a candidate change.
- Use standing checks instead of internal clone.
- Spot what improved what held and what weakened.
- Choose Choose, Hold, or Iterate for a clear reason.
- Close with monitoring follow-up, not false certainty.

How to get the most from a mission

- Read the headline before you read the candidate result.
- Bridge noisy checks first to the important trade offs stay stable.
- Use the interpretation help when a check looks ambiguous.
- Use the debrief to compare your path with the strongest prospect path.

Optional extension

Some missions include a second external model comparison step. That step is still part of the workflow workflow because it lets you compare outside output against the same standing checks and occlusion discipline.

Mission Catalog

Filter Bar

Right-side
Orientation Panels

Key Insight: Start foundational. The Support Summary Compression Update is the ideal first run.

Callout: Use filters only when they help narrow the scenario list; they are optional.

Strategic Mission Routing

	Internal-Only	External Step
Core Watchtower Pack	Cleanest first lessons in baseline protection and release judgment. Contained within the app's workspace.	Requires manually copying packets to an external LLM (e.g., ChatGPT) and pasting back.
Optional Challenge Pack	Pushes harder on governance and safety boundaries; polished candidates may be unsafe. Contained within the app's workspace.	Advanced external comparison testing strict discipline against live, un-sandboxed outputs.


Phase 2: Analyze the Evidence Foundation

The Rule


Read the incident and proposed change first.

The Anchor

The baseline is the standard for what must remain visible or enforceable. The candidate result is evaluated against this standard, not in isolation.



Phase 3: Classify Protected Behaviors



High-Risk: These checks reveal critical trade-offs and safety boundaries. Read these first.

Low-Risk: Polish wins and readability improvements that can distract you into approving for the wrong reasons.

Core Strategy:
Prioritize Risky Checks.

The Classification Engine

Improved (+)

Improved: Made the protected behavior stronger without undermining controls.

Held

Held: Preserved the behavior the baseline was already protecting.

Tradeoff

Tradeoff: Improved something, but at a meaningful operational cost.

Weakened (-)

Weakened: Made an important baseline protection less visible, enforceable, or reliable.

Course Correcting Mid-Flight

Study Cues:

Use Interpretation Help to reveal study cues if a check feels ambiguous.

Current Signals:

Higher bars mean stronger fit—except for Regression Risk, where a higher bar indicates more danger.

Study cue

Protects in deremation: Candidate preserves the term only only wthor compact imanges.

Study cue: Your read matches the strongest practice interpretation.

Improved

Held

Tradeoff

Weakened

Expected read: Held.

► Interpretation help

CURRENT SIGNALS

See what moved before you judge

Evidær bars crontly mean stronger fit—for regression risk, a higher bar means more danger, not more inoneer.

Evidence Adherence Strong



Evidence references are presented clearly.

Format Stability Strong



The new fielded structure semjiated intact.

Refusal Correctness Held



No risky loos or comp-pasted advice approved.

Regression Risk Low




The compact version is archved the key protected insolution.

Phase 4: Choose a Proportionate Action



Release: Protections preserved, improvement bounded.



Hold: Release-blocking weakness found (safety, boundary, refusal).



Iterate: Has promise, but needs revision.




Monitoring Follow-Up

The Stewardship Rule

Monitoring is not a fix for a release blocker. It manages bounded residual uncertainty after a proportionate decision is made.

Phase 5: Calibrate via Matrix Comparison



Delta Summary: See exactly where your read aligned or drifted.

- **Matrix Comparison:**
Compare your path with the strongest-practice path.

The Goal: Replay the reasoning pattern until you can explain the decision without teaching aids.

Handling Manual External Comparisons



1. Copy Packet



2. Run in External LLM



3. Copy Output



4. Paste to Workspace


The Mechanism

The app does not run an LLM internally.

The Purpose

This optional step tests your ability to apply strict standing-check discipline to live, un-sandboxed outputs instead of trusting novelty or polish alone.

Watchtower Core Vocabulary



Golden Rules for the User

- ✓ Protect the baseline before rewarding polish.
- ✓ Read risky checks first.
- ✓ Treat Release, Hold, and Iterate as proportionate steps, not praise labels.
- Use monitoring for bounded uncertainty, not not to compensate for weakness.

Quick Reference State Machine

Route	What to do	What to watch for
Launcher	Choose a mission	Start foundational before moving to safety scenarios.
Packet & Checks	Read with baseline-first discipline	High-risk trade-offs carry more weight than polish.
Signals & Decision	Lock a proportionate choice	Regression risk runs opposite to positive fit.
Debrief	Study the matrix and strongest-practice path	Explain the pattern in your own words.